

# HW5: Clustering

Savannah

2025-10-21

## Libraries Needed:

```
library(tidyverse)
library(cluster)
library(glue)
library(gridExtra)
```

#Task 1: Clusters and the Gap Statistic

## Data Loading:

From the data\_generation.py file, I generated the necessary datasets according to the directions on the PDF. I will now proceed to load all datasets.

```
file_list <- list.files("derived_data_1", pattern = "*.csv", full.names = TRUE)

# Read all CSVs into a named list of data frames
df_list <- file_list %>%
  set_names(~ basename(.) %>% tools::file_path_sans_ext()) %>%
  map(read_csv)

n <- c(2:6)
side_length <- c(1,10,2:9)
```

## Clustergap Simulations:

For each data set, I will run the cluster gap function. Then, I will visualize the results via a plot of estimated clusters by side\_length.

## Plotting Estimated Number of Clusters by Side Length

To create the estimated number of clusters by side\_length plot, I will first divide df\_list into lists of datasets by their values of n. Then, I will run the clusGap function on each individual list, saving the optimal number of clusters for each dataset. Lastly, I will generate five plots, one per n value.

```
D2 <- df_list[grepl("^df_2_", names(df_list))]
D3 <- df_list[grepl("^df_3_", names(df_list))]
D4 <- df_list[grepl("^df_4_", names(df_list))]
D5 <- df_list[grepl("^df_5_", names(df_list))]
D6 <- df_list[grepl("^df_6_", names(df_list))]
```

```
estimate_clusters <- function(df_list, n_value) {

  results <- list()
```

```

for (name in names(df_list)) {
  df <- df_list[[name]]

  # Drop label column if it exists
  if ("label" %in% names(df)) {
    df <- df %>% select(-label)
  }

  # Extract side_length from name (e.g., df_2_5 + 5)
  side_length <- as.integer(str_extract(name, "\\d+"))

  # Run clusGap (you can increase B later)
  set.seed(123)

  gap <- clusGap(df, FUN = kmeans, K.max = n_value + 5, nstart = 20, iter.max = 50, B= 10)

  # Get optimal number of clusters using firstSEmax
  opt_nc <- maxSE(gap$Tab[, "gap"], gap$Tab[, "SE.sim"], method = "firstSEmax")

  # Store results
  results[[name]] <- tibble(
    dataset = name,
    n = n_value,
    side_length = side_length,
    optimal_num_clusters = opt_nc
  )
}

# Combine and return
bind_rows(results)
}

gap_D2 <- estimate_clusters(D2, n_value = 2)
gap_D3 <- estimate_clusters(D3, n_value = 3)
gap_D4 <- estimate_clusters(D4, n_value = 4)
gap_D5 <- estimate_clusters(D5, n_value = 5)
gap_D6 <- estimate_clusters(D6, n_value = 6)

```

Now, I will first plot each gap statistic, and then I want to combine all of these ggplots into one figure for simplicity.

```

# Create individual plots
p2 <- ggplot(gap_D2, aes(x = side_length, y = optimal_num_clusters)) +
  geom_point() + geom_line() + ggtitle("n = 2") +
  geom_hline(yintercept = 2, linetype = "dashed", color = "red")

p3 <- ggplot(gap_D3, aes(x = side_length, y = optimal_num_clusters)) +
  geom_point() + geom_line() + ggtitle("n = 3") +
  geom_hline(yintercept = 3, linetype = "dashed", color = "red")

p4 <- ggplot(gap_D4, aes(x = side_length, y = optimal_num_clusters)) +
  geom_point() + geom_line() + ggtitle("n = 4") +
  geom_hline(yintercept = 4, linetype = "dashed", color = "red")

```


```

p5 <- ggplot(gap_D5, aes(x = side_length, y = optimal_num_clusters)) +
  geom_point() + geom_line() + ggtitle("n = 5") +
  geom_hline(yintercept = 5, linetype = "dashed", color = "red")

p6 <- ggplot(gap_D6, aes(x = side_length, y = optimal_num_clusters)) +
  geom_point() + geom_line() + ggtitle("n = 6") +
  geom_hline(yintercept = 6, linetype = "dashed", color = "red")

grid.arrange(p2, p3, p4, p5, p6, ncol = 3)

```



From these plots, we find that for  $n = 2$ , the optimal number of clusters is predicted is reduced after  $\text{side\_length} = 1$ . For  $n = 3$  and  $4$ , the optimal number of clusters is reduced for  $\text{side\_length} < 3$ . At  $n = 5$ ,  $\text{side\_length} < 4$  is when the optimal number of clusters is reduced. Under  $n = 6$ ,  $\text{side\_length} < 5$  is when the estimated number of clusters drastically reduces.

To save it separately as a figure:

```

if (!dir.exists("figures_1")) {
  dir.create("figures_1")
}

png("figures_1/cluster_gap_grid.png", width = 1000, height = 800)

grid.arrange(p2, p3, p4, p5, p6, ncol = 2)

invisible(dev.off())

```

## Task 2: Spectral Clustering on Concentric Shells

This section will continue the work on task 2. The file `data_generation_2.py` has the beginning of this task, with the function for creating data, a test of the data generation function, an interactive plot of this sample data, and the creation as well as saving of the data sets with max radii from 10 to 0.

Within this section, I will first implement the function to transform the data such that the clusters are differentiable with kmeans. Then, I will implement the `clusGap` function to see at which point clusters are more challenging to distinguish. Lastly, I will generate a plot of the estimated number of clusters by the maximum radius allowed.

### Data Loading:

```
file_list <- list.files("derived_data_2", pattern = "*.csv", full.names = TRUE)

# Read all CSVs into a named list of data frames
df_list2 <- file_list %>%
  set_names(~ basename(.) %>% tools::file_path_sans_ext()) %>%
  map(read_csv)
```

### Spectral Clustering Transformation Function

```
# nested functions:

# Adjacency Matrix Creation -

adjmatrix <- function(data, d_threshold) {
  data <- as.matrix(data)

  dist_matrix <- as.matrix(dist(data, method = "euclidean"))
  adjacency_matrix <- ifelse(dist_matrix < d_threshold, 1, 0)

  diag(adjacency_matrix) <- 0 #doesn't make sense for a node to have a euclidean distance with itself

  return(adjacency_matrix)
}

# Normalized Laplacian Function -

Laplacian <- function(adj_matrix) {

  adj_matrix <- as.matrix(adj_matrix)
  degree_vector <- rowSums(adj_matrix)
  D <- diag(degree_vector)
  L <- D - adj_matrix

  # avoid division by zero
  inv_sqrt_deg <- 1 / sqrt(degree_vector)
  inv_sqrt_deg[!is.finite(inv_sqrt_deg)] <- 0 # set Inf or NaN to 0

  D_inv_sqrt <- diag(inv_sqrt_deg)

  # symmetric normalized Laplacian
  L_sym <- D_inv_sqrt %*% L %*% D_inv_sqrt
```

```

    return(L_sym)
}

# Eigendecomposition and Extracting the K Eigenvectors -

EigenVectors <- function(matrix_input, k) {

  # Ensure input is numeric matrix
  matrix_input <- as.matrix(matrix_input)

  # Check if matrix is square
  if (nrow(matrix_input) != ncol(matrix_input)) {
    stop("Matrix must be square.")
  }

  eig <- eigen(matrix_input)

  idx <- order(eig$values)
  eigenvectors <- eig$vectors[, idx[1:k], drop = FALSE]

  # Return as numeric matrix
  return(eigenvectors)
}

```

Putting it all together:

```

# Full Spectral Clustering Transformation Function -

Spectral <- function(x,k){
  # x is the dataset, k is the number of clusters (in our case, 4)
  A <- adjmatrix(x, 1)
  L <- Laplacian(A)
  V_set <- as.matrix(EigenVectors(L, k))

  return(list(cluster = kmeans(V_set, k)$cluster))
}

set.seed(123)
results <- list()


for (name in names(df_list2)) {
  df <- df_list2[[name]]
  gap <- clusGap(df, FUN = Spectral, K.max = 10, B= 10)
  opt_nc <- maxSE(gap$Tab[, "gap"], gap$Tab[, "SE.sim"], method = "firstSEmax")
  results[[name]] <- tibble(
    dataset = name,
    max_radius = as.integer(sub("df__", "", name)),
    optimal_num_clusters = opt_nc
  )
}

results2 <- bind_rows(results)

```

## Plotting the Results:

```
ggplot(results2, aes(x = max_radius, y = optimal_num_clusters)) +  
  geom_point() + geom_line() + ggtitle("n = 4") +  
  geom_hline(yintercept = 4, linetype = "dashed", color = "red")
```



```
if (!dir.exists("figures_2")) {  
  dir.create("figures_2")  
}  
  
png("figures_2/Spectral_Cluster_Max_Radius.png", width = 1000, height = 800)  
  
ggplot(results2, aes(x = max_radius, y = optimal_num_clusters)) +  
  geom_point() + geom_line() + ggtitle("n = 4") +  
  geom_hline(yintercept = 4, linetype = "dashed", color = "red")  
  
invisible(dev.off())
```

While the algorithm does not perform all that well, it seems as though the performance is particularly poor upon reaching a max\_radius of 3 and below. I believe the algorithm would perform better if the d\_threshold was decreased potentially, as that would force the identification of more tight-knit clusters and worse if the d\_threshold were increased (would result in harder identification of clusters and more room for error).